

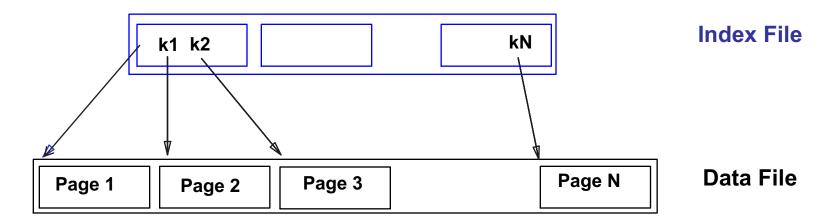
B+-tree

#### Outline

- B+-tree
- Query
- Update
  - Insertion
  - Deletion

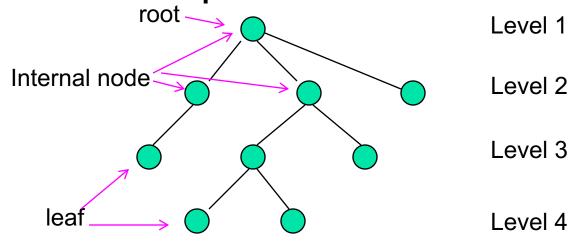


- Find all students with gpa > 3.0
  - If records are sorted on gpa, do binary search to find first such student, then scan to find others.
  - Cost of binary search can be quite high.
- Simple idea: Create an "index" file.



## B+ Tree: Most Widely Used Index

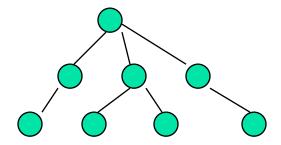
General concept of a tree:



- I. Height of a node: its distance to the root
- II. If a higher level node is connected to a lower level node, then the higher level node is called a parent (grandparent, ancestor, etc.) of the lower level node

### B+ Tree (cont.)

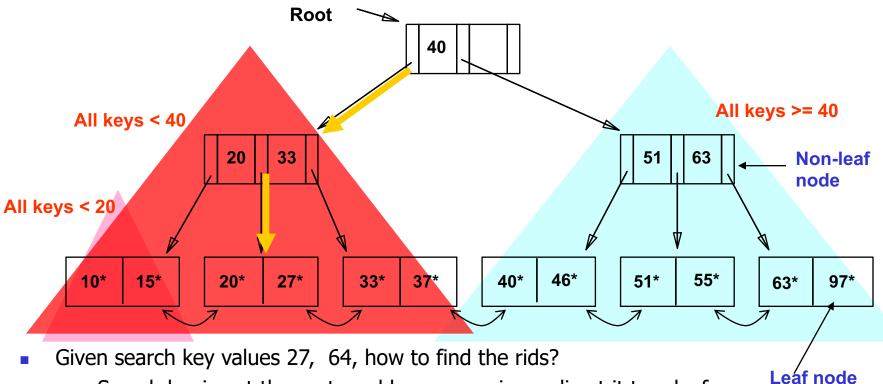
- Balanced tree: all leaves at the same level
  - Example:



- Structure of a B+ tree:
  - It is balanced: all leaf nodes are at the same level
  - Each node has a special structure

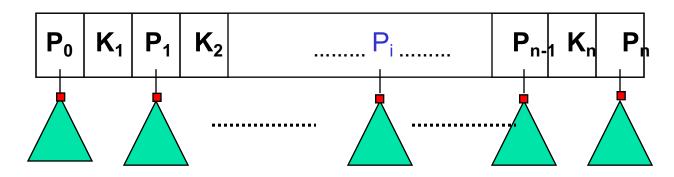
#### Example: A B+ tree

In this example, each node holds at least 1 entry, and at most 2 entries



- Search begins at the root, and key comparisons direct it to a leaf
- Note that key values are sorted at each level **COMP231**

## B+ tree: Internal node structure



Each P<sub>i</sub> is a pointer to a child node, each K<sub>i</sub> is a search key value Pointers outnumber search key values by *exactly one*.

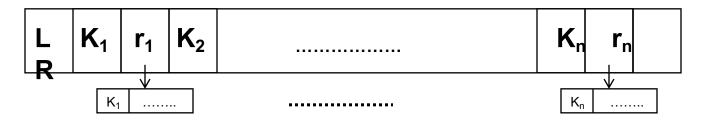


#### Requirements:

- K1 < K2 < ... < Kn
- If the node is not the root, we require d ≤ n ≤ 2d where d is a pre-determined value for this B+ tree, called its *order*
- If the node is the root, we require  $1 \le n \le 2d$
- For any search key value K in the subtree pointed by Pi,
  - If Pi = P0, we require that K < K1
  - If Pi = P1, ..., Pn-1, we require that  $Ki \le K < Ki+1$
  - If Pi = Pn, we require Kn ≤ K

### 4

#### B+ tree: leaf node structure



- Each r<sub>i</sub> is a pointer to a record that contains search key value K<sub>i</sub>
- L points to the left neighbor, and R points to the right neighbor
- $K_1 < K_2 < ... < K_n$
- We require  $d \le n \le 2d$  where d is the order of this B+ tree

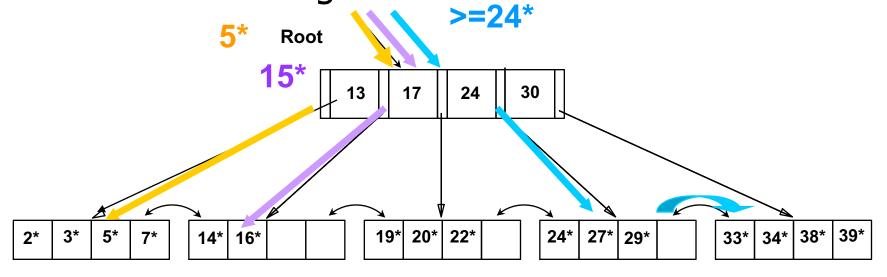
### Outline

- B+-tree
- Query
- Update
  - Insertion
  - Deletion

### Example: a B+ tree with order

2

- Search for 5\*, 15\*, all data entries >= 24\* ...
- The last one is a range search, we need to do the sequential scan, starting from the first leaf containing a value >= 24.



COMP231



#### Height/Depth of B+-tree

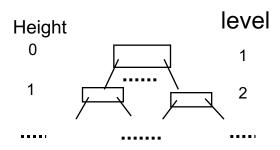
What is the minimum height of the B+tree?

What is the maximum height of the B+tree?

#### Minimum height of the B+tree

- In general nodes are pages
- Let H be the height of the B+ tree: need to read H+1 pages to reach a leaf node
- Let F be the number of pointers in a node (for internal node, called *fanout*)
  - Level 1 = 1 page  $= F^0$  page
  - Level 2 = F pages = F¹ pages
  - Level 3 = Fx F pages = F<sup>2</sup> pages
  - Level H+1 = ..... = F<sup>H</sup> pages (i.e., leaf nodes)
  - Suppose there are D data entries. So there are at least D/(F-1) leaf nodes

    Note that data entry is one less
  - D/(F-1) <= F<sup>H</sup>. That is, H >=  $log_F(\frac{D}{F-1})$



than pointers

F = 2d + 1

#### Maximum height of the B+tree

What is the maximum height of the B+-tree?

COMP231



#### Cost of Searching a key

- Usually, each node occupies a page
- Cost of searching a key = H + 1 (pages)

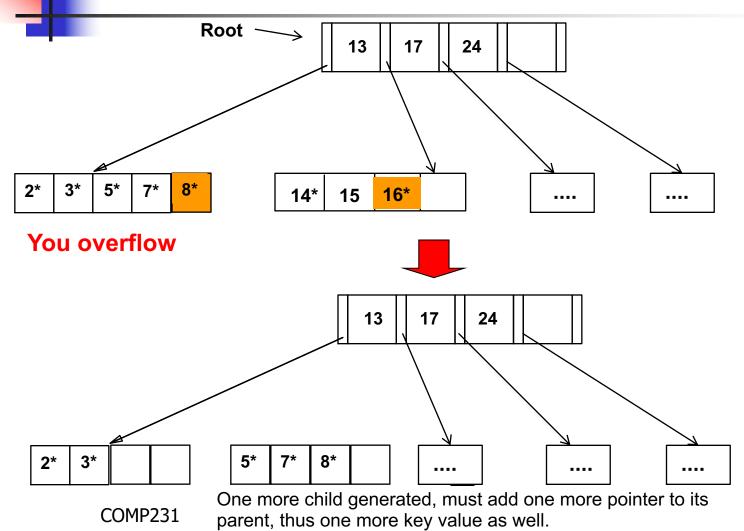
COMP231

### Outline

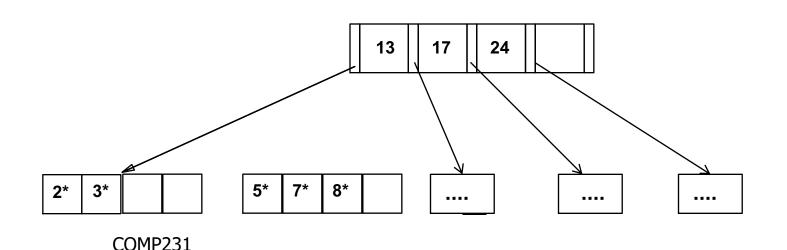
- B+-tree
- Query
- Update
  - Insertion
  - Deletion

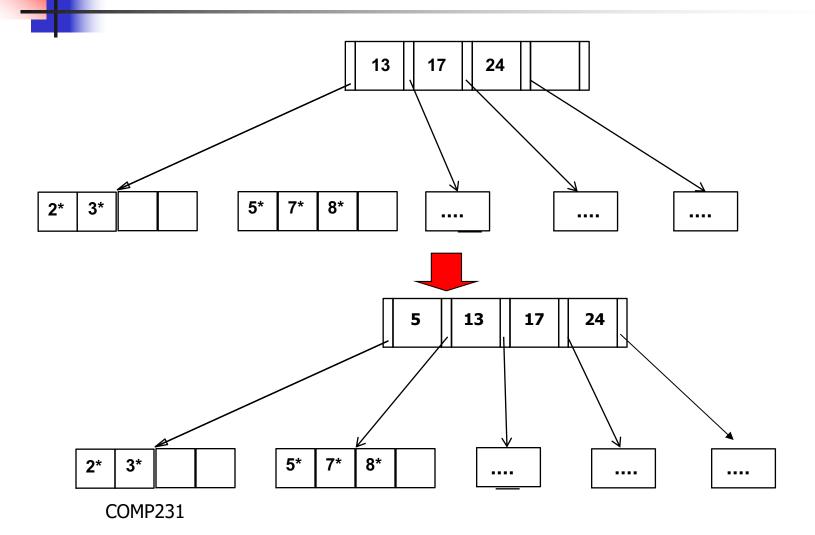
#### Inserting a Data Entry into a B+ Tree

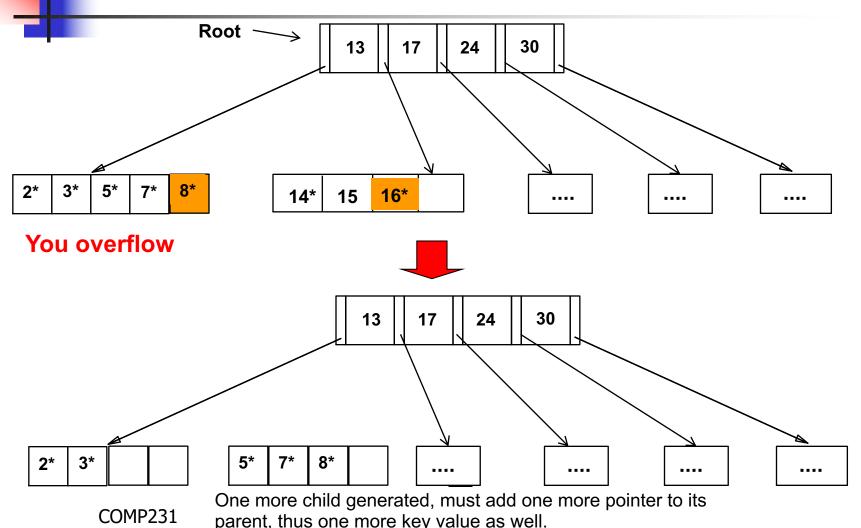
- Find correct leaf L.
- Put data entry onto L.
  - If L has enough space, done!
  - Else, must *split L (into L and a new node L2)* 
    - Redistribute entries evenly, put middle key in L2
    - **copy up** middle key.
    - Insert index entry pointing to L2 into parent of L.
- This can happen recursively
  - To split index node, redistribute entries evenly,
  - but <u>push up</u> middle key. (Contrast with leaf splits.)
- Splits "grow" tree; root split increases height.
  - Tree growth: gets <u>wider</u> or <u>one level taller at top.</u>



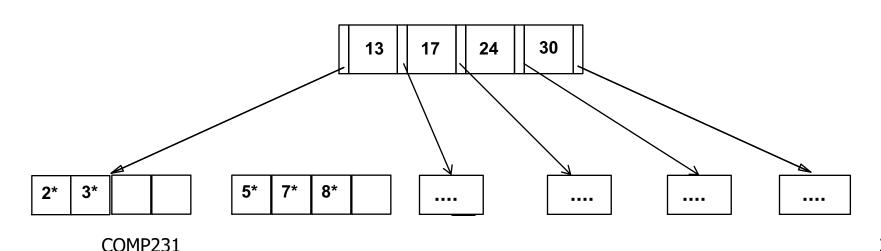


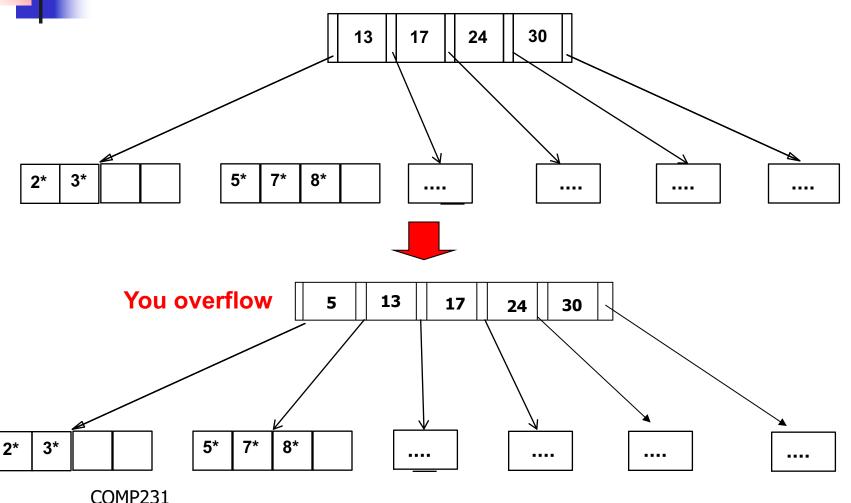




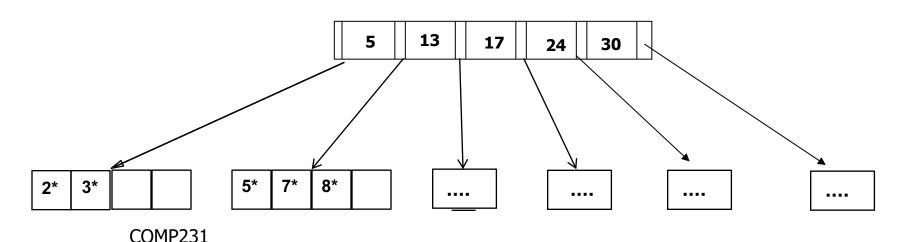








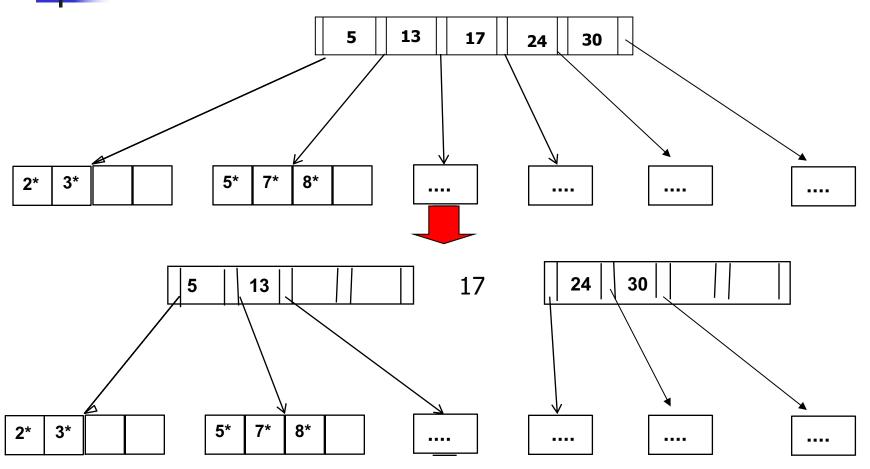






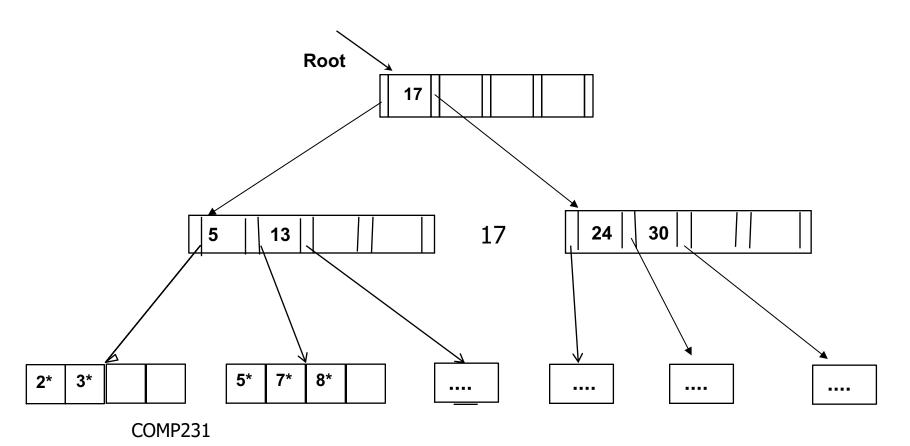
COMP231

## Example 2: Inserting 16\*, 8\* into Example B+ tree



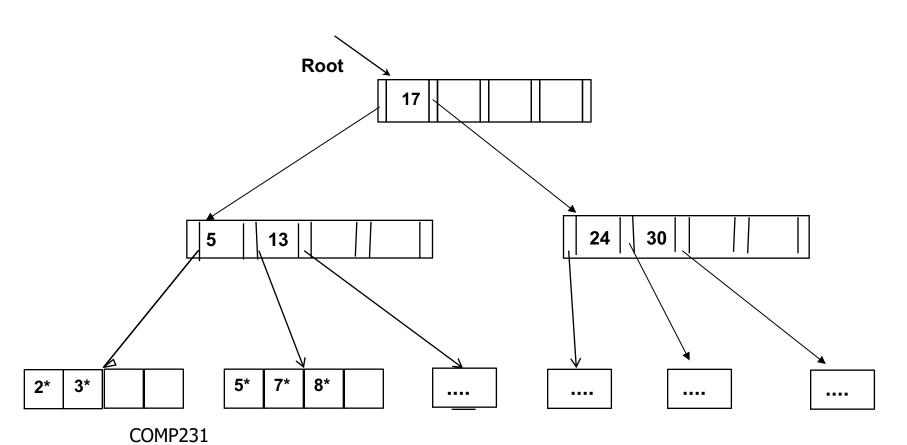
25





26





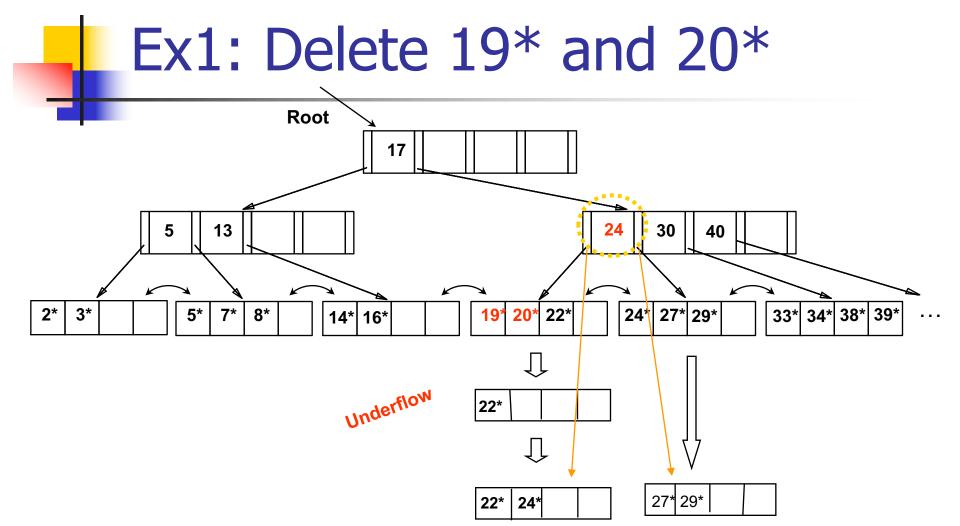
### Outline

- B+-tree
- Query
- Update
  - Insertion
  - Deletion

#### Deleting a Data Entry from a B+ Tree

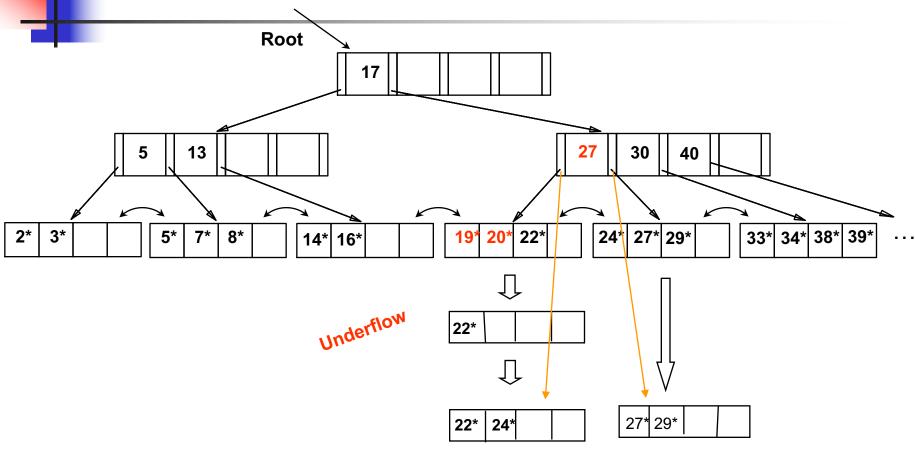
- Start at root, find leaf L where entry belongs.
- Remove the entry.
  - If L is at least half-full, done!
  - If L has only d-1 entries,
    - Try to re-distribute, borrowing from <u>sibling</u> (adjacent node with same parent as L).
    - If re-distribution fails, <u>merge</u> L and sibling.
- If merge occurred, must delete entry (pointing to L or sibling) from parent of L.
- Merge could propagate to root, decreasing height.

COMP231



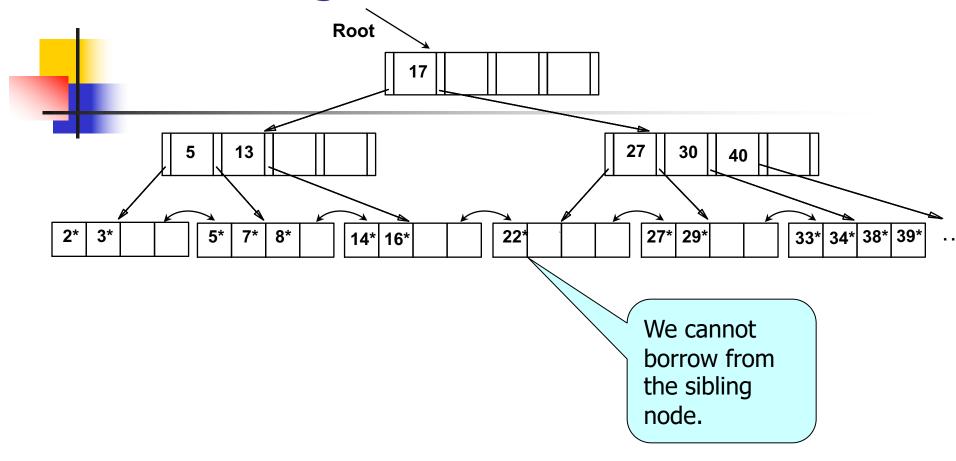
COMP231 30

# Ex1: Delete 19\* and 20\*



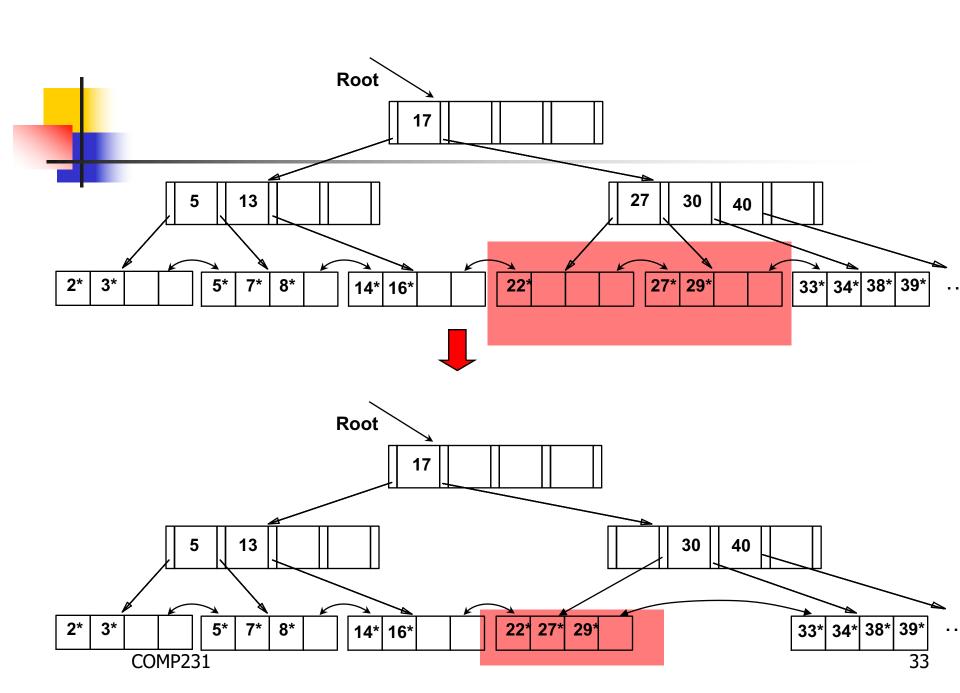
COMP231

#### Ex2: Deleting 24\*

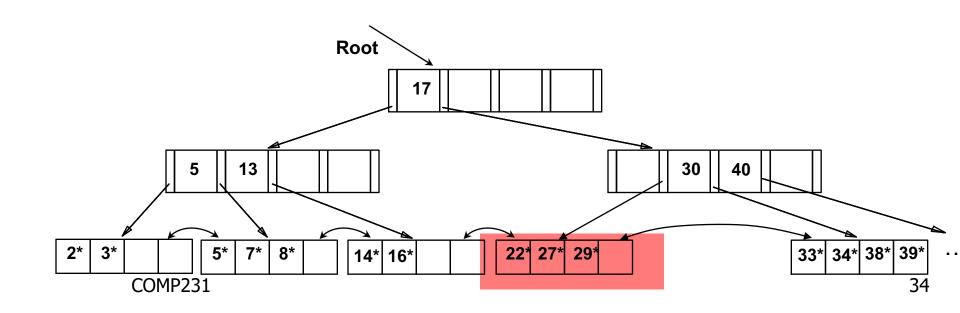


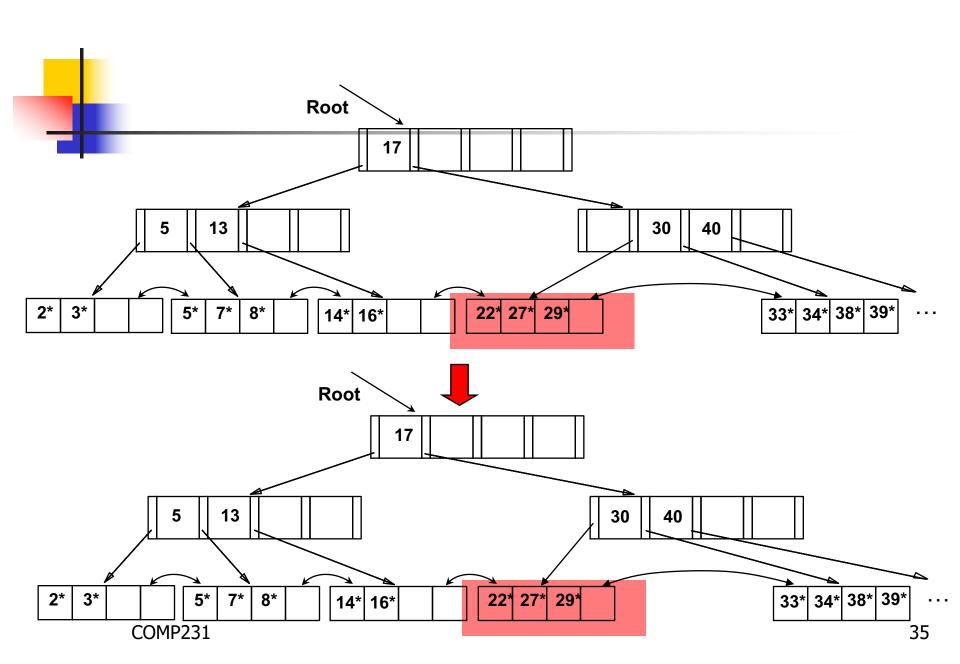
COMP231

32

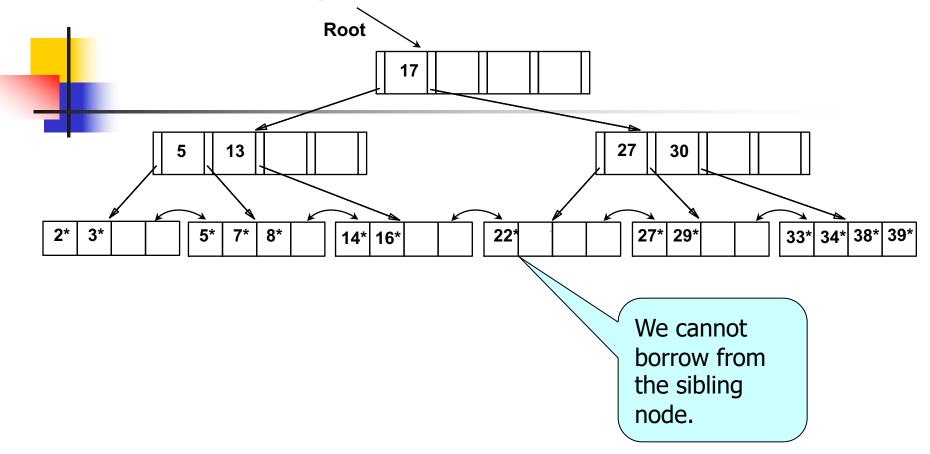








#### Ex3: Deleting 24\*



COMP231 36

