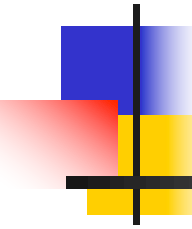

COMP231

Join Algorithm

Prepared by Raymond Wong

Presented by Raymond Wong

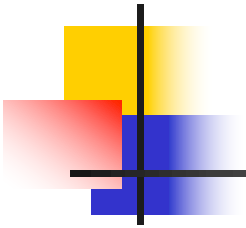

Example

- Consider the following SQL query.
  - E.g., student (sid, sname, age)
    take (sid, cid)

```
select sname
from student S, take T
where S.sid = T.sid
```

Student: 500 pages, 80 tuples/page, 50 bytes/tuple
Take: 1000 pages, 100 tuples/page, 40 bytes/tuple

- 200 courses in table Take
- 1, 2, …, 40 in attribute Age of table Student

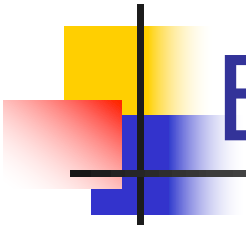

Example

- Assuming the following table sizes:
  - Student: 500 pages, 80 tuples/page, 50 bytes/tuple
  - Take: 1000 pages, 100 tuples/page, 40 bytes/tuple
- Assume that there are 200 courses in table Take

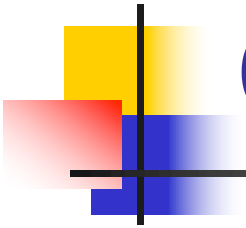

Outline

- Simple-Nested Loop Join
- Block-Nested Loop Join
- Sort-Merge Join
- Index-Nested Loop Join
- Hash Join

Student: 500 pages, 80 tuples/page, 50 bytes/tuple
Take: 1000 pages, 100 tuples/page, 40 bytes/tuple

- 200 courses in table Take
- 1, 2, …, 40 in attribute Age of table Student

Simple-Nested Loop Join

For each page p of T (T is called ***outer relation***)
  for each page q of S (S is called ***inner relation***)
    for each tuple t ∈ p and each tuple s ∈ q
      such that t.sid = s.sid
      output <t, s> to the output

For each tuple r in page p
  For each tuple s in page q,
    If r.sid = s.sid then
      output the tuple joined from
      tuple r and tuple s

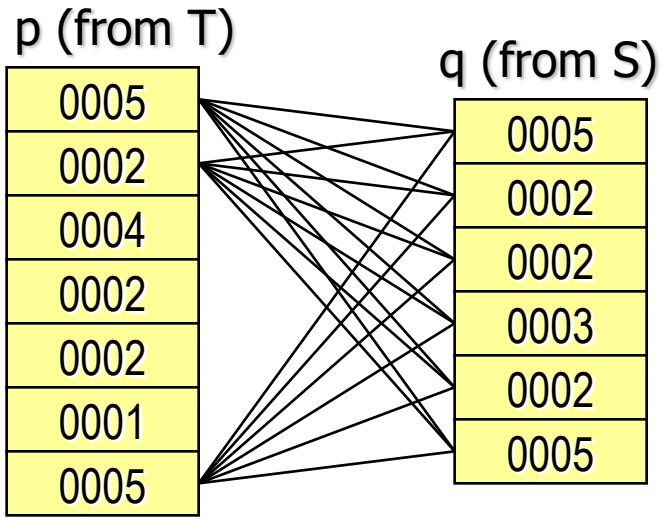

Simple-Nested Loop Join

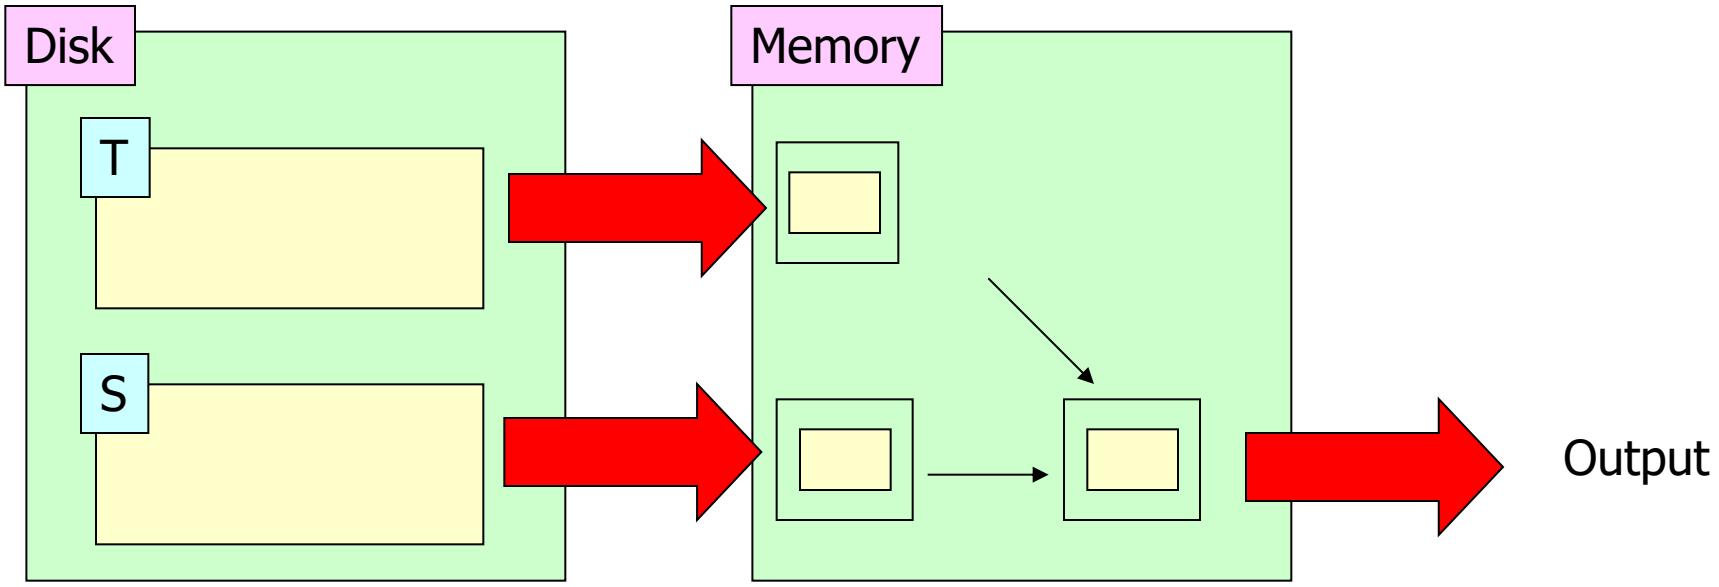

Student: 500 pages, 80 tuples/page, 50 bytes/tuple
Take: 1000 pages, 100 tuples/page, 40 bytes/tuple

- 200 courses in table Take
- 1, 2, …, 40 in attribute Age of table Student

Simple-Nested Loop Join

For each page p of T (T is called ***outer relation***)
  for each page q of S (S is called ***inner relation***)
    for each tuple t $\in$ p and each tuple s $\in$ q
      such that t.sid = s.sid
      output <t, s> to the output

We need 3 pages for buffer

1 page for p (from T)

1 page for q (from S)

1 page for the output

Student: 500 pages, 80 tuples/page, 50 bytes/tuple
Take: 1000 pages, 100 tuples/page, 40 bytes/tuple

- 200 courses in table Take
- 1, 2, …, 40 in attribute Age of table Student

Simple-Nested Loop Join

For each page p of T (T is called ***outer relation***)
  for each page q of S (S is called ***inner relation***)
    for each tuple t ∈ p and each tuple s ∈ q
      such that t.sid = s.sid
      output <t, s> to the output

Cost of Reading T = 1000 pages

The total number of times that S is read = 1000

Cost of Reading S (with multiple times) = 1000*500 = 500000 pages

Total Cost = 1000 + 500000 = 501000 pages

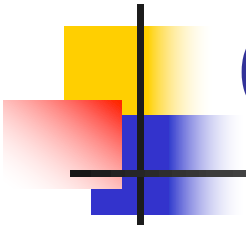

Outline

- Simple-Nested Loop Join
- Block-Nested Loop Join
- Sort-Merge Join
- Index-Nested Loop Join
- Hash Join

Student: 500 pages, 80 tuples/page, 50 bytes/tuple
Take: 1000 pages, 100 tuples/page, 40 bytes/tuple

- 200 courses in table Take
- 1, 2, …, 40 in attribute Age of table Student

Block-Nested Loop Join

For each block P of T (T is called ***outer relation***)
  for each page q of S (S is called ***inner relation***)
    for each tuple t ∈ P and each tuple s ∈ q
      such that r.sid = s.sid
      output <t, s> to the output

Block-Nested Loop Join

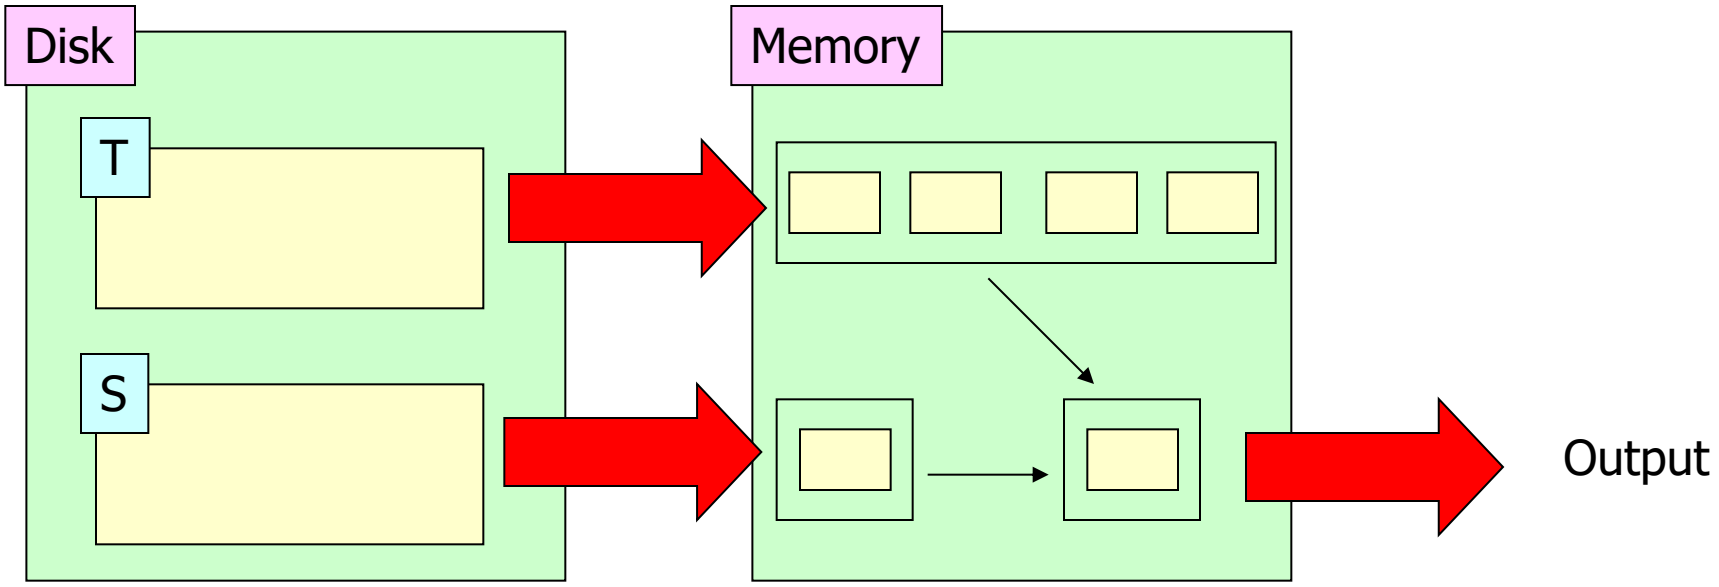

Student: 500 pages, 80 tuples/page, 50 bytes/tuple
Take: 1000 pages, 100 tuples/page, 40 bytes/tuple

- 200 courses in table Take
- 1, 2, …, 40 in attribute Age of table Student

Block-Nested Loop Join

Suppose that B = 6 pages

For each block P of T (T is called ***outer relation***)
  for each page q of S (S is called ***inner relation***)
    for each tuple t ∈ P and each tuple s ∈ q
      such that t.sid = s.sid
      output <t, s> to the output

Suppose we have B pages in memory (or B buffer pages)

B-2 pages for P (from T)

1 page for q (from S)

1 page for the output

Student: 500 pages, 80 tuples/page, 50 bytes/tuple
Take: 1000 pages, 100 tuples/page, 40 bytes/tuple

- 200 courses in table Take
- 1, 2, …, 40 in attribute Age of table Student

Block-Nested Loop Join

Suppose that B = 6 pages

For each block P of T (T is called ***outer relation***)
  for each page q of S (S is called ***inner relation***)
    for each tuple t ∈ P and each tuple s ∈ q
      such that t.sid = s.sid
      output <t, s> to the output

Cost of Reading T = 1000 pages

The total number of times that S is read = ⌈1000/(6-2)⌉ = 250

Cost of Reading S (with multiple times) = 250*500 = 125000 pages

Total Cost = 1000 + 125000 = 126000 pages

Student: 500 pages, 80 tuples/page, 50 bytes/tuple
Take: 1000 pages, 100 tuples/page, 40 bytes/tuple

- 200 courses in table Take
- 1, 2, …, 40 in attribute Age of table Student

Block-Nested Loop Join

Suppose that B = 6 pages

- T is the outer relation
- S is the inner relation
- How about the following?
  - T is the inner relation
  - S is the outer relation

Student: 500 pages, 80 tuples/page, 50 bytes/tuple
Take: 1000 pages, 100 tuples/page, 40 bytes/tuple

- 200 courses in table Take
- 1, 2, …, 40 in attribute Age of table Student

Block-Nested Loop Join

Suppose that B = 6 pages

For each block P of S (S is called ***outer relation***)
  for each page q of T (T is called ***inner relation***)
    for each tuple s ∈ P and each tuple t ∈ q
      such that t.sid = s.sid
      output <t, s> to the output

Cost of Reading S = 500 pages

The total number of times that T is read = ⌈500/(6-2)⌉ = 125

Cost of Reading T (with multiple times) = 125*1000 = 125000 pages

Total Cost = 500 + 125000 = 125500 pages

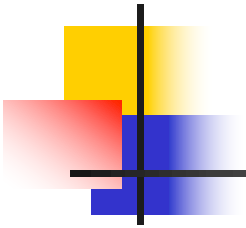

Block-Nested Loop Join

- Execution 1
  - T is the outer relation
  - S is the inner relation
  - Cost = 126000 pages
- Execution 2
  - S is the inner relation
  - T is the outer relation
  - Cost = 125500 pages
- Which one is better?

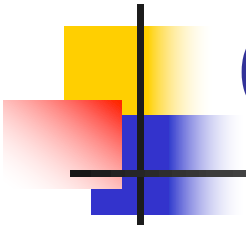

Outline

- Simple-Nested Loop Join
- Block-Nested Loop Join
- Sort-Merge Join
- Index-Nested Loop Join
- Hash Join

Sort-Merge Join

1. Sort table T according to sid
2. Sort table S according to sid
3. Merge table T and table S

T

0004
0005
0001
0002
0002
0005
0002

S

0003
0002
0005
0002
0005
0002

Sort-Merge Join

1. Sort table T according to sid
2. Sort table S according to sid
3. Merge table T and table S

T

0001
0002
0002
0002
0004
0005
0005

S

0003
0002
0005
0002
0005
0002

Sort-Merge Join

1. Sort table T according to sid
2. Sort table S according to sid
3. Merge table T and table S

T

0001
0002
0002
0002
0004
0005
0005

S

0002
0002
0002
0003
0005
0005

Sort-Merge Join

1. Sort table T according to sid
2. Sort table S according to sid
3. Merge table T and table S

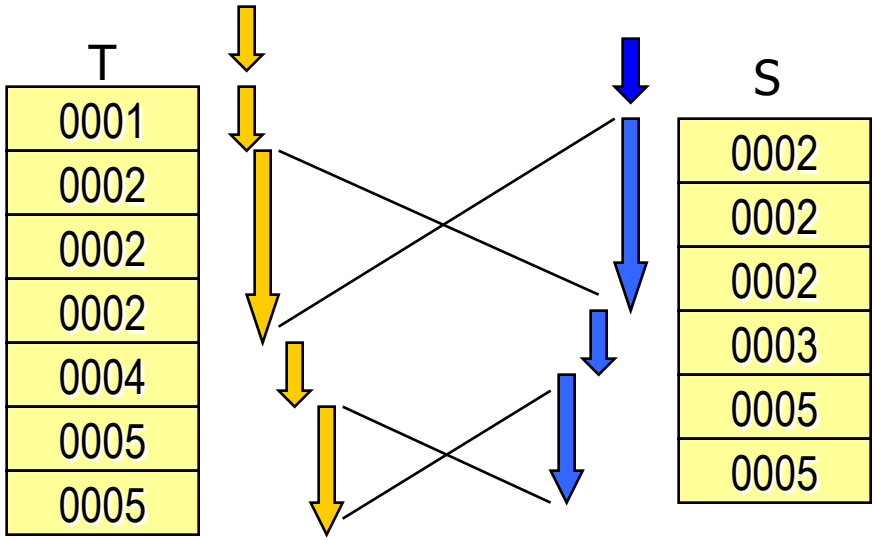

Sort-Merge Join

1. Sort table T according to sid
2. Sort table S according to sid
3. Merge table T and table S

Cost of Sorting = 2N * (# of passes)

where no. of passes= $1+\lceil \log_{B-1}\lceil N/B \rceil \rceil$

Student: 500 pages, 80 tuples/page, 50 bytes/tuple
Take: 1000 pages, 100 tuples/page, 40 bytes/tuple

- 200 courses in table Take
- 1, 2, …, 40 in attribute Age of table Student

Sort-Merge Join

Suppose that B = 6 pages

1. Sort table T according to sid
2. Sort table S according to sid
3. Merge table T and table S

Cost of Sorting T = $2*1000*(1+\lceil \log_{6-1} \lceil 1000/6 \rceil \rceil)$
= 10000

Cost of Sorting S = $2*500*(1+\lceil \log_{6-1} \lceil 500/6 \rceil \rceil)$
= 4000

Cost of Merging = Cost of Reading T + Cost of Reading S
= 1000 + 500
= 1500

Total Cost = 10000 + 4000 + 1500 = 15500 pages

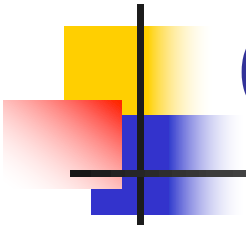

Outline

- Simple-Nested Loop Join
- Block-Nested Loop Join
- Sort-Merge Join
- Index-Nested Loop Join
- Hash Join

Student: 500 pages, 80 tuples/page, 50 bytes/tuple
Take: 1000 pages, 100 tuples/page, 40 bytes/tuple

- 200 courses in table Take
- 1, 2, …, 40 in attribute Age of table Student

Index-Nested Loop Join

- Assume that there is an index built on attribute sid for table Student
- E.g., If we search for the tuples with sid = 0005

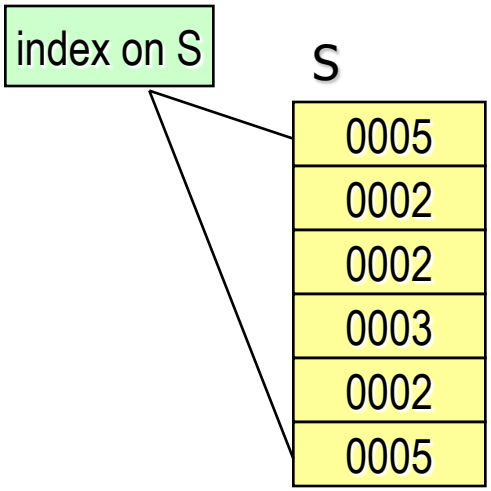

Index-Nested Loop Join

For each tuple t of T
    ID ← t.sid
    use the index to obtain a tuple s (or tuples) in S
        with s.sid = ID
    for each such s in S
        output <t, s> to the output

Assume that we have a hash index on sid for table S
The cost of accessing a hash index is 1.2 pages

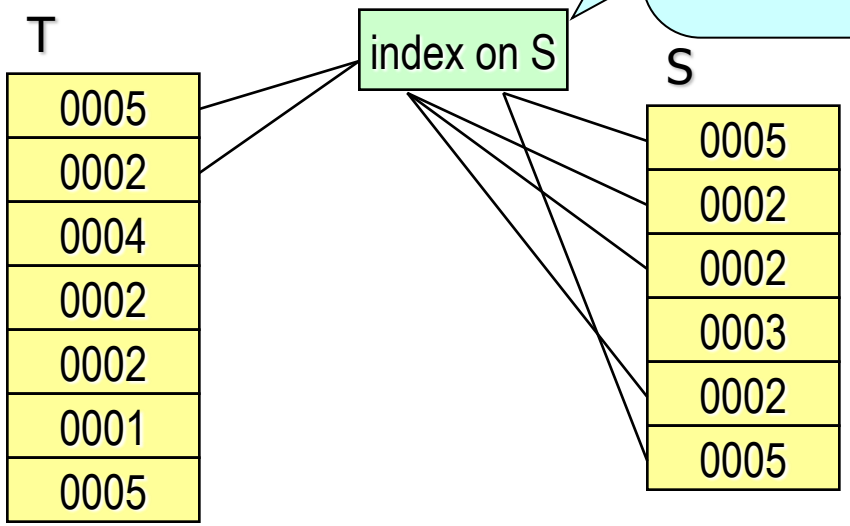

Student: 500 pages, 80 tuples/page, 50 bytes/tuple
Take: 1000 pages, 100 tuples/page, 40 bytes/tuple

- 200 courses in table Take
- 1, 2, …, 40 in attribute Age of table Student

Index-Nested Loop Join

Suppose that B = 6 pages

For each tuple t of T
    ID ← t.sid
    use the index to obtain a tuple s (or tuples) in S
        with s.sid = ID
    for each such s in S
        output <t, s> to the output

Cost of Reading T = 1000 pages

The total number of tuples in T =1000*100=100000

Why "+1"?

Cost of Reading S (with multiple times) =100000*(1.2 + 1)= 220000 pages

Total Cost = 1000 + 220000 = 221000 pages

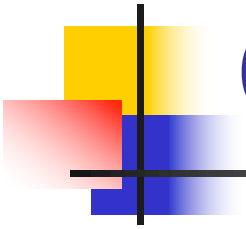

Outline

- Simple-Nested Loop Join
- Block-Nested Loop Join
- Sort-Merge Join
- Index-Nested Loop Join
- Hash Join

Hash Join

1. Hash all tuples in T based on sid
2. Hash all tuples in S based on sid
3. For each hash slot,
    1. Join all tuples t from T and all tuples s from S where t.sid = s.sid
    2. Output all these tuples

Hash Join

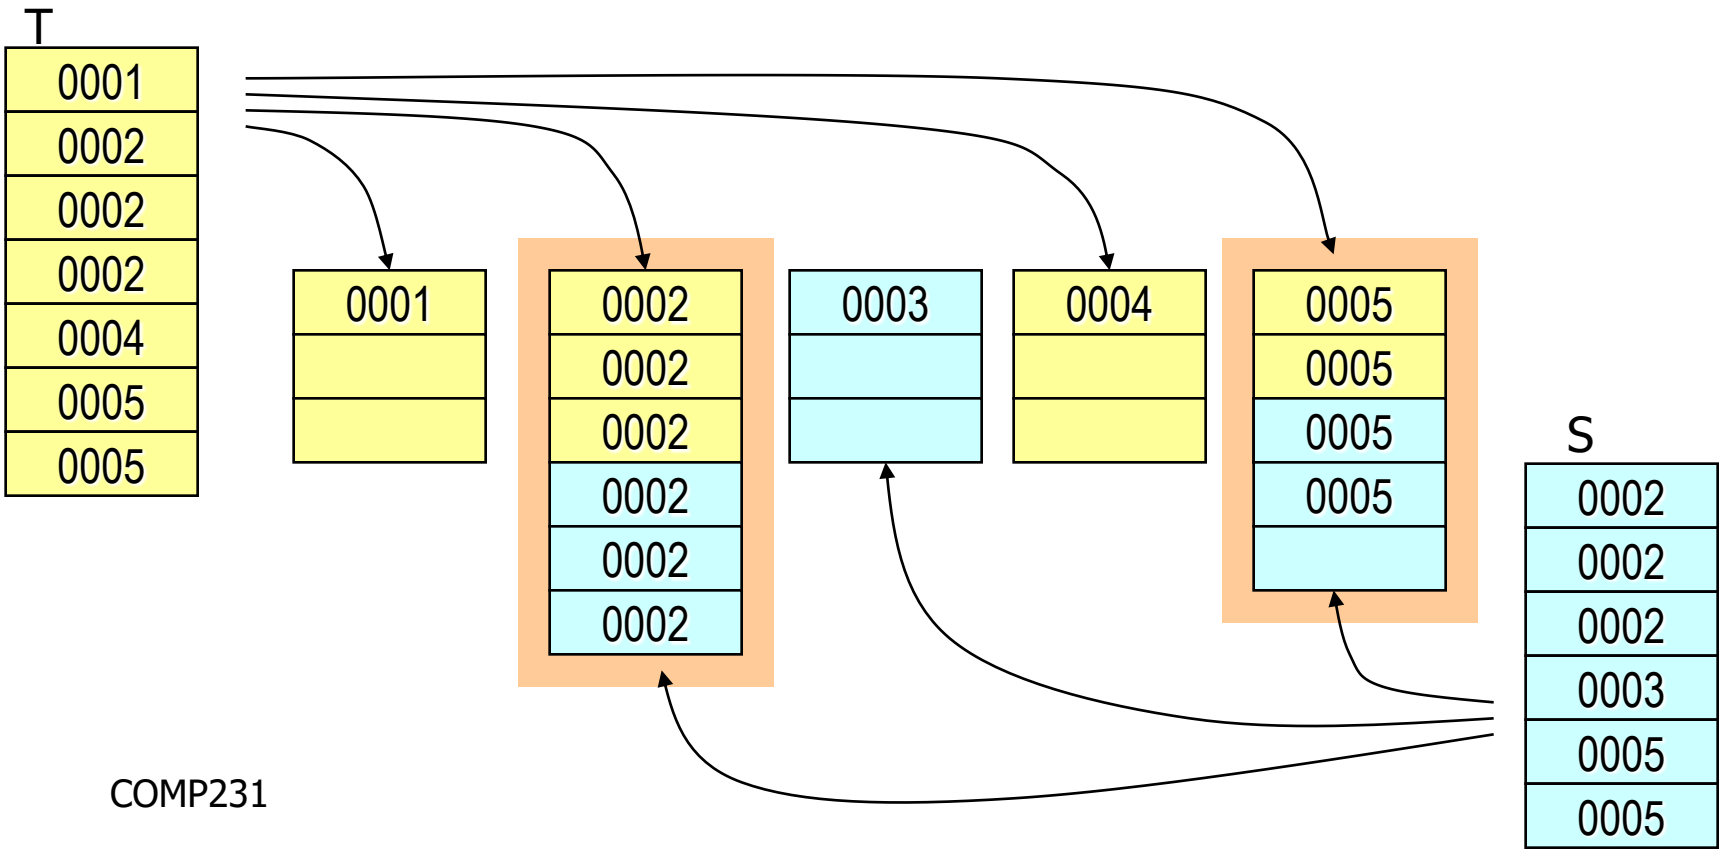
T
0001
0002
0002
0002
0004
0005
0005
0001
0002
0002
0002
0002
0002
0003
0004
0005
0005
0005
0005
S
0002
0002
0002
0003
0005
0005